

Spotify song popularity analysis

Completed as part of the coursework for the Multivariate Analysis and Statistical Modelling module at the University of Brighton, 2025/26.

Gosia Szkatula

2026-06-10

Table of contents

1	Spotify	2
1.1	Executive summary	2
1.2	Introduction	2
1.3	Data preparation	2
1.4	Part 1 - Global models	7
1.4.1	Multiple linear regression models	7
1.4.2	Variable selection	8
1.4.3	Main effects	9
1.4.4	Diagnostics	11
1.5	Part 2 - Genre-specific models	12
1.5.1	Model estimation	12
1.5.2	Cross validation	16
1.6	Conclusions	17
1.6.1	Global model	17
1.6.2	Genre-specific models	17
1.6.3	Limitations	18

List of Figures

1	Spotify. Plot of the main effects coefficients	10
2	Spotify. Plot of F-statistics from ANOVA table	11
3	Spotify. Diagnostic plots	12
4	Spotify. Bar plots of coefficients by genre and audio variable	16

List of Tables

1	Spotify. Comparison of Global Linear Models	7
3	Spotify. ANOVA table	11
4	Spotify. CV summary table	17

```
library(tidyverse)
library(rio)
library(olsrr)
library(broom)
library(blorr)
library(forecast)
```

1 Spotify

Data source:

<https://www.kaggle.com/datasets/melissamonfared/spotify-tracks-attributes-and-popularity>

1.1 Executive summary

This analysis examines how audio features influence song popularity on Spotify. Results show that more danceable and louder tracks tend to be more popular, whereas features such as instrumentality and energy are associated with lower popularity.

Genre plays an important role. Some relationships, such as the positive effect of danceability and loudness, were consistent across genres, while others, including energy and duration, vary by music type. This suggests that popularity depends not only on audio characteristics but also on genre-specific dynamics.

Overall, audio features provide useful insight but explain only part of popularity, indicating that external factors such as artist recognition or marketing likely play a significant role.

1.2 Introduction

The Spotify Tracks Dataset [<https://www.kaggle.com/datasets/melissamonfared/spotify-tracks-attributes-and-popularity>] contains 114,000 observations across 21 variables, including track, artist, genre, popularity, and audio features derived from the Spotify API. The dataset appears artificially balanced, with roughly 1,000 observations per genre, although the sampling process is not documented.

The primary objective is to identify which audio variables best explain variation in popularity. A secondary objective is to assess whether these relationships vary by genre.

1.3 Data preparation

Data were inspected and variables grouped into:

- information variables (excluded),
- categorical variables (inspected but not used),
- audio variables (numeric predictors),
- target variable (popularity)

Observations with zero popularity were removed. Genres were grouped into broader categories (Pop, HipHop, Rock, Metal, Electronic, Classical) based on key terms, excluding unmatched tracks. Tracks may appear multiple times with different genre labels; therefore, observations were collapsed to one row per track to ensure independence and avoid duplication. Where multiple genre assignments existed, a single genre was retained, introducing a simplifying assumption that may not fully capture the multi-genre nature of some tracks. Only the target, genre, and numeric audio variables were retained, with numeric variables standardised.

The final sample size was $n = 28,970$ observations.

```
spotify <- import("data/dataset.csv")

# Create factors where appropriate
spotify$track_genre <- as.factor(spotify$track_genre)
spotify$key <- as.factor(spotify$key)
spotify$mode <- as.factor(spotify$mode)
spotify$time_signature <- as.factor(spotify$time_signature)

# Datatypes
str(spotify)
```

```
'data.frame': 114000 obs. of 21 variables:
 $ index      : int  0 1 2 3 4 5 6 7 8 9 ...
 $ track_id   : chr   "5Su0ikwiRyPMVoIQDJUgSV" "4qPNDBW1i3p13qLCt0Ki3A" "1iJBSr7s
 $ artists    : chr   "Gen Hoshino" "Ben Woodward" "Ingrid Michaelson;ZAYN" "Kina
 $ album_name : chr   "Comedy" "Ghost (Acoustic)" "To Begin Again" "Crazy Rich As
 $ track_name : chr   "Comedy" "Ghost - Acoustic" "To Begin Again" "Can't Help Fa
 $ popularity : int   73 55 57 71 82 58 74 80 74 56 ...
 $ duration_ms : int   230666 149610 210826 201933 198853 214240 229400 242946 189
 $ explicit   : logi   FALSE FALSE FALSE FALSE FALSE FALSE ...
 $ danceability : num   0.676 0.42 0.438 0.266 0.618 0.688 0.407 0.703 0.625 0.442
 $ energy      : num   0.461 0.166 0.359 0.0596 0.443 0.481 0.147 0.444 0.414 0.63
 $ key         : Factor w/ 12 levels "0","1","2","3",...: 2 2 1 1 3 7 3 12 1 2 ...
 $ loudness    : num   -6.75 -17.23 -9.73 -18.52 -9.68 ...
 $ mode        : Factor w/ 2 levels "0","1": 1 2 2 2 2 2 2 2 2 2 ...
 $ speechiness : num   0.143 0.0763 0.0557 0.0363 0.0526 0.105 0.0355 0.0417 0.036
 $ acousticness : num   0.0322 0.924 0.21 0.905 0.469 0.289 0.857 0.559 0.294 0.426
 $ instrumentalness: num   1.01e-06 5.56e-06 0.00 7.07e-05 0.00 0.00 2.89e-
06 0.00 0.00 4.19e-03 ...
 $ liveness    : num   0.358 0.101 0.117 0.132 0.0829 0.189 0.0913 0.0973 0.151 0.
 $ valence     : num   0.715 0.267 0.12 0.143 0.167 0.666 0.0765 0.712 0.669 0.196
 $ tempo       : num   87.9 77.5 76.3 181.7 119.9 ...
 $ time_signature : Factor w/ 5 levels "0","1","3","4",...: 4 4 4 3 4 4 3 4 4 4 ...
 $ track_genre : Factor w/ 114 levels "acoustic","afrobeat",...: 1 1 1 1 1 1 1 1 1
```

```
# create variable groups for easier manipulation
info_vars <- c(
  "index", "track_id", "artists",
  "album_name", "track_name"
)

other_vars <- c(
  "explicit", "key", "mode",
  "time_signature"
)

audio_vars <- c(
  "duration_ms", "danceability", "energy",
  "loudness", "speechiness", "acousticness",
  "instrumentalness", "liveness",
  "valence", "tempo"
)

target <- "popularity"
```

```
# Distribution
summary(spotify[, other_vars])
```

explicit	key	mode	time_signature
Mode :logical	7	:13245	0: 163
FALSE:104253	0	:13061	1: 973
TRUE :9747	2	:11644	3: 9195
	9	:11313	4:101843
	1	:10772	5: 1826
	5	: 9368	
	(Other):44597		

```
summary(spotify[, audio_vars])
```

duration_ms	danceability	energy	loudness
Min. : 0	Min. :0.0000	Min. :0.0000	Min. : -49.531
1st Qu.: 174066	1st Qu.:0.4560	1st Qu.:0.4720	1st Qu.: -10.013
Median : 212906	Median :0.5800	Median :0.6850	Median : -7.004
Mean : 228029	Mean :0.5668	Mean :0.6414	Mean : -8.259
3rd Qu.: 261506	3rd Qu.:0.6950	3rd Qu.:0.8540	3rd Qu.: -5.003
Max. :5237295	Max. :0.9850	Max. :1.0000	Max. : 4.532

speechiness	acousticness	instrumentalness	liveness
Min. :0.00000	Min. :0.0000	Min. :0.00e+00	Min. :0.0000
1st Qu.:0.03590	1st Qu.:0.0169	1st Qu.:0.00e+00	1st Qu.:0.0980
Median :0.04890	Median :0.1690	Median :4.16e-05	Median :0.1320
Mean :0.08465	Mean :0.3149	Mean :1.56e-01	Mean :0.2136
3rd Qu.:0.08450	3rd Qu.:0.5980	3rd Qu.:4.90e-02	3rd Qu.:0.2730
Max. :0.96500	Max. :0.9960	Max. :1.00e+00	Max. :1.0000

valence	tempo
Min. :0.0000	Min. : 0.00
1st Qu.:0.2600	1st Qu.: 99.22
Median :0.4640	Median :122.02
Mean :0.4741	Mean :122.15
3rd Qu.:0.6830	3rd Qu.:140.07
Max. :0.9950	Max. :243.37

```
summary(spotify[, target])
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.00	17.00	35.00	33.24	50.00	100.00

```
spotify <- spotify %>%
```

```
# 1. Map genres
```

```
mutate(
  genre = case_when(
    track_genre %in% c("pop", "k-pop", "j-pop", "mandopop", "cantopop",
                      "indie-pop", "power-pop", "synth-pop", "pop-film") ~ "Pop",

    track_genre %in% c("hip-hop", "r-n-b") ~ "HipHop",

    track_genre %in% c("rock", "alt-rock", "rock-n-roll", "rockabilly",
                      "hard-rock", "psych-rock", "alternative", "indie",
                      "grunge", "emo", "j-rock", "punk-rock") ~ "Rock",

    track_genre %in% c("metal", "heavy-metal", "death-metal",
                      "black-metal", "metalcore") ~ "Metal",

    track_genre %in% c("electronic", "electro", "edm",
                      "house", "techno", "trance",
                      "dubstep", "drum-and-bass", "breakbeat", "trip-hop") ~ "Elect",

    track_genre %in% c("classical", "opera") ~ "Classical",

    TRUE ~ NA_character_
  )
) %>%
```

```
# 2. Keep rows with genre identified
```

```
filter(!is.na(genre), popularity != 0) %>%
```

```
# 3. Collapse to one row per track (to fix duplication)
```

```
group_by(track_id) %>%
```

```
slice_head(n = 1) %>%
```

```
ungroup() %>%
```

```

# Further processing
dplyr::select(-all_of(c(info_vars, other_vars))) %>%
mutate(across(all_of(audio_vars), scale)) %>%
dplyr::select(-track_genre)

spotify %>%
  count(track_id) %>%
  filter(n > 1)

rio::export(spotify, "outputs/spotify_clean.xlsx")

# genre - factor
spotify$genre <- as.factor(spotify$genre)

# Datatypes - check again
str(spotify)

tibble [28,970 x 12] (S3: tbl_df/tbl/data.frame)
 $ popularity      : int [1:28970] 62 38 20 86 36 58 36 16 44 33 ...
 $ duration_ms     : num [1:28970, 1] -0.462 -0.582 0.11 -0.16 -0.55 ...
 ..- attr(*, "scaled:center")= num 240386
 ..- attr(*, "scaled:scale")= num 108513
 $ danceability    : num [1:28970, 1] 0.848 0.0808 -0.7724 -1.159 -
0.9074 ...
 ..- attr(*, "scaled:center")= num 0.541
 ..- attr(*, "scaled:scale")= num 0.163
 $ energy          : num [1:28970, 1] 0.31 0.965 0.822 0.922 0.349 ...
 ..- attr(*, "scaled:center")= num 0.698
 ..- attr(*, "scaled:scale")= num 0.231
 $ loudness        : num [1:28970, 1] 0.941 0.679 1.189 0.519 -0.197 ...
 ..- attr(*, "scaled:center")= num -7.31
 ..- attr(*, "scaled:scale")= num 4.01
 $ speechiness     : num [1:28970, 1] 1.7406 0.0512 -0.0568 0.0349 -
0.5361 ...
 ..- attr(*, "scaled:center")= num 0.0723
 ..- attr(*, "scaled:scale")= num 0.0676
 $ acousticness    : num [1:28970, 1] -0.567 -0.697 -0.745 -0.759 -
0.753 ...
 ..- attr(*, "scaled:center")= num 0.227
 ..- attr(*, "scaled:scale")= num 0.298
 $ instrumentalness: num [1:28970, 1] -0.496 -0.1871 -0.489 -0.496 0.0685 ...
 ..- attr(*, "scaled:center")= num 0.141
 ..- attr(*, "scaled:scale")= num 0.285
 $ liveness        : num [1:28970, 1] -0.726 0.688 -0.132 -0.628 0.424 ...
 ..- attr(*, "scaled:center")= num 0.209
 ..- attr(*, "scaled:scale")= num 0.174
 $ valence         : num [1:28970, 1] 1.59 1.029 -1.207 -0.844 0.322 ...

```

```

..- attr(*, "scaled:center")= num 0.445
..- attr(*, "scaled:scale")= num 0.248
$ tempo      : num [1:28970, 1] 1.22 1.955 1.666 0.759 -1.248 ...
..- attr(*, "scaled:center")= num 125
..- attr(*, "scaled:scale")= num 29.7
$ genre      : Factor w/ 6 levels "Classical","Electronic",...: 3 6 2 6 6 6 6 2 ...

# final sample size
nrow(spotify)

[1] 28970

```

1.4 Part 1 - Global models

1.4.1 Multiple linear regression models

Three linear models were fitted:

- `lm_no_genre`: baseline without genre, only numeric audio variables
- `lm_main`: includes genre
- `lm_interaction`: includes genre interactions

Model performance improves with complexity, with adjusted R^2 increasing from approximately 0.09 (no genre) to 0.14 (with genre) and 0.18 (interaction model), (Table 1). This indicates that genre contributes meaningfully and that relationships vary across genres, even as overall explanatory power was limited.

```

data = spotify

# fit linear models
lm_no_genre <- lm(popularity ~ . - genre, data = data)
lm_main <- lm(popularity ~ ., data = data)
lm_interaction <- lm(popularity ~ . * genre, data = data)

no_genre <- t(glance(lm_no_genre))
main <- t(glance(lm_main))
interaction <- t(glance(lm_interaction))

model_table <- cbind(no_genre, main, interaction)

colnames(model_table) <- c("No Genre", "With Genre", "Interaction")

model_table |>
  round(3) |>
  knitr::kable() |>
  kableExtra::kable_styling(bootstrap_options = c("striped", "hover", "condensed"))

```

Table 1: Spotify. Comparison of Global Linear Models

	No Genre	With Genre	Interaction
r.squared	0.094	0.138	0.184
adj.r.squared	0.094	0.138	0.182
sigma	18.408	17.951	17.489
statistic	299.937	310.266	100.010
p.value	0.000	0.000	0.000
df	10.000	15.000	65.000
logLik	-125485.076	-124753.546	-123974.156
AIC	250994.153	249541.092	248082.311
BIC	251093.441	249681.750	248636.670
deviance	9813249.789	9329961.418	8841213.601
df.residual	28959.000	28954.000	28904.000
nobs	28970.000	28970.000	28970.000

1.4.2 Variable selection

Variable selection methods (`ols_step_backward`, `step`, and `drop1`), retained all predictors except speechiness, and the model was subsequently refitted.

```
lm_no_genre <- lm(popularity ~ . - genre, data = data)
lm_no_genre_no_speechiness <- lm(popularity ~ . - genre - speechiness, data = data)

model <- lm_no_genre_no_speechiness

# ols p-value method
model_ols <- {
  capture.output(tmp <- ols_step_backward_p(model))
  tmp
}

# AIC method
model_step <- step(model, direction = "backward", trace = 0)

# drop method
model_drop <- {
  capture.output(tmp <- drop1(model, test = "F"))
  tmp
}

# Extract variables
attr(terms(model_ols$model), "term.labels")

[1] "duration_ms"      "danceability"      "energy"            "loudness"
[5] "acousticness"     "instrumentalness"  "liveness"          "valence"
[9] "tempo"
```



```
attr(terms(model_step), "term.labels")
```

```
[1] "duration_ms"      "danceability"      "energy"             "loudness"
[5] "acousticness"     "instrumentalness"  "liveness"           "valence"
[9] "tempo"
```

```
model_drop[!is.na(model_drop[, "Pr(>F)"]) &
  model_drop[, "Pr(>F)"] < 0.05, ]
```

	Df	Sum of Sq	RSS	AIC	F value	Pr(>F)
duration_ms	1	5853.683	9819437	168793.1	17.274287	0.0000324
danceability	1	199464.286	10013048	169358.8	588.621445	0.0000000
energy	1	100524.122	9914108	169071.1	296.647862	0.0000000
loudness	1	133021.750	9946605	169165.9	392.548743	0.0000000
acousticness	1	5770.813	9819354	168792.9	17.029738	0.0000369
instrumentalness	1	213296.224	10026880	169398.8	629.439657	0.0000000
liveness	1	35076.387	9848660	168879.2	103.510829	0.0000000
valence	1	21361.511	9834945	168838.8	63.038069	0.0000000
tempo	1	3127.167	9816711	168785.1	9.228306	0.0023851

1.4.3 Main effects

To ensure interpretability, analysis focuses on the no genre model, using numerical predictors only. The coefficient plot (Figure 1) shows that loudness and danceability have strong positive effects, whereas energy and instrumentalness have strong negative effects. Other variables have weaker negative effects.

```
tidy(model, conf.int = TRUE) %>%
  filter(term != "(Intercept)") %>%
  filter(term != "genre") %>%
  ggplot(aes(x = estimate, y = reorder(term, estimate))) +
  geom_point(colour = "steelblue") +
  geom_errorbarh(aes(xmin = conf.low, xmax = conf.high), height = 0.3) +
  geom_vline(xintercept = 0, linetype = "dashed", colour = "red") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 7)) +
  labs(title = "Coefficient Plot (Main Effects Model)",
    x = "Estimate", y = NULL)
```




Figure 1: Spotify. Plot of the main effects coefficients

ANOVA results (Table 3, Figure 2) confirm this, with loudness and danceability having the largest F statistics, indicating the greatest explanatory power. Acousticness and tempo contribute minimally relative to other variables.

```
tidy(anova(model)) %>%
  filter(term != "Residuals") %>%
  ggplot(aes(x = statistic,
             y = reorder(term, statistic))) +
  geom_col(fill = "steelblue") +
  labs(title = "F-statistics from ANOVA table",
       x = "F statistic",
       y = NULL) +
  theme_bw()
```




Figure 2: Spotify. Plot of F-statistics from ANOVA table

```
anova_model <- anova(model) |>
  knitr::kable(digits = 3)

anova_model
```

Table 3: Spotify. ANOVA table

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
duration_ms	1	43297.708	43297.708	127.772	0.000
danceability	1	343742.095	343742.095	1014.387	0.000
energy	1	30102.292	30102.292	88.832	0.000
loudness	1	346982.692	346982.692	1023.950	0.000
acousticness	1	5004.453	5004.453	14.768	0.000
instrumentalness	1	187026.426	187026.426	551.917	0.000
liveness	1	34047.486	34047.486	100.475	0.000
valence	1	22723.042	22723.042	67.056	0.000
tempo	1	3127.167	3127.167	9.228	0.002
Residuals	28960	9813583.546	338.867	NA	NA

1.4.4 Diagnostics

Diagnostics (Figure 3) indicate violations of normality and homoscedasticity. However, given the large sample size, these do not materially affect estimates, and the model is retained for interpretation rather than strict inference.

```
par(mfrow = c(2, 2))
plot(model)
```




Figure 3: Spotify. Diagnostic plots

1.5 Part 2 - Genre-specific models

1.5.1 Model estimation

The data were split into genre-specific subsets, and separate models fitted. Variable selection using `drop1()` was applied, and models were refitted using retained predictors.

Genre-specific models reveal both consistent and varying relationships (Figure 4). Danceability is positively associated with popularity across all genres (except not retained in HipHop), while liveness and valence (excl. Metal) were consistently negative. Loudness is positive in most genres but negative for Classical. These results align broadly with the no genre model (Figure 1).

In contrast, energy, instrumentality, and duration vary in both magnitude and direction, indicating strong genre dependence and important heterogeneity.

```
# genre sample size (n)
summary(spotify$genre)
```

Classical	Electronic	HipHop	Metal	Pop	Rock
1421	7508	1402	3927	7213	7499

```
# split into subsets
groups <- split(spotify, spotify$genre)

for (i in seq_along(groups)) {
  assign(paste0("genre_", i), groups[[i]])
}

# name subsets, drop genre
Classical <- dplyr::select(groups[[1]], -genre)
Electronic <- dplyr::select(groups[[2]], -genre)
HipHop <- dplyr::select(groups[[3]], -genre)
Metal <- dplyr::select(groups[[4]], -genre)
Pop <- dplyr::select(groups[[5]], -genre)
Rock <- dplyr::select(groups[[6]], -genre)

# check subset dimensions
dim(Classical) # Classical (n=1421)

[1] 1421  11

dim(Electronic) # Electronic (n=7508)

[1] 7508  11

dim(HipHop) # HipHop (n=1402)

[1] 1402  11

dim(Metal) # Metal (n=3927)

[1] 3927  11

dim(Pop) # Pop (n=7213)

[1] 7213  11

dim(Rock) # Rock (n=7499)

[1] 7499  11
```

```

run_genre_model <- function(data) {

  # 1. Fit full model (no genre inside subset)
  model_full <- lm(popularity ~ ., data = data)

  # 2. Variable selection using drop1
  drop_tbl <- drop1(model_full, test = "F")

  # Convert to data frame for easier handling
  drop_df <- as.data.frame(drop_tbl)

  # Keep variables that are significant
  keep_vars <- rownames(drop_df)[
    !is.na(drop_df$`Pr(>F)`) & drop_df$`Pr(>F)` < 0.05
  ]

  # Remove "<none>" if present
  keep_vars <- setdiff(keep_vars, "<none>")

  # 3. Refit model with selected variables
  formula_new <- as.formula(
    paste("popularity ~", paste(keep_vars, collapse = " + "))
  )

  model_final <- lm(formula_new, data = data)

  # 4. Return everything useful
  list(
    full_model = model_full,
    selected_vars = keep_vars,
    final_model = model_final
  )
}

```

```

res_classical <- run_genre_model(Classical)
res_pop <- run_genre_model(Pop)
res_rock <- run_genre_model(Rock)
res_metal <- run_genre_model(Metal)
res_hiphop <- run_genre_model(HipHop)
res_electronic <- run_genre_model(Electronic)

# extract coefficients from each model
extract_coefs <- function(res, genre_name) {
  coefs <- summary(res$final_model)$coefficients
  df <- as.data.frame(coefs)
  df$Variable <- rownames(df)
  df$Genre <- genre_name
}

```

```

    df
  }

# combine
coef_all <- rbind(
  extract_coefs(res_classical, "Classical"),
  extract_coefs(res_pop, "Pop"),
  extract_coefs(res_rock, "Rock"),
  extract_coefs(res_metal, "Metal"),
  extract_coefs(res_hiphop, "HipHop"),
  extract_coefs(res_electronic, "Electronic")
)

# plot
ggplot(coef_all[coef_all$Variable != "(Intercept)", ],
  aes(x = Genre, y = Estimate, fill = Genre)) +
  geom_col() +
  facet_wrap(~ Variable, scales = "free_y") +
  geom_hline(yintercept = 0, linetype = "dashed") +
  theme_bw() +
  theme(
    legend.position = "none",
    axis.text.x = element_text(angle = 90, vjust = 0.5, hjust = 1)
  ) +
  labs(title = "Variation in Coefficients by Genre",
    x = NULL, y = "Estimate")

```




Figure 4: Spotify. Bar plots of coefficients by genre and audio variable

1.5.2 Cross validation

Cross-validation evaluates predictive performance (Table 4.), where lower values indicate better generalisation. Although genre-specific models have lower R^2 , some achieve lower CV errors, suggesting improved prediction within genres. This highlights a trade-off between explanatory power and predictive accuracy.

```
compare_models_cv <- function(...) {
  models <- list(...)

  map_dfr(
    models,
    CV,
    .id = "model"
  ) |>
```



```

knitr::kable(digits = 3)
}

compare_models_cv(
  No_Genre = lm_no_genre,
  No_Genre_Speechiness = lm_no_genre_no_speechiness,
  With_Genre = lm_main,
  Genre_Interaction = lm_interaction,
  Pop = res_pop$final_model,
  Rock = res_rock$final_model,
  HipHop = res_hiphop$final_model,
  Metal = res_metal$final_model,
  Electronic = res_electronic$final_model,
  Classical = res_classical$final_model
)

```

Table 4: Spotify. CV summary table

model	CV	AIC	AICc	BIC	AdjR2
No_Genre	338.983	168780.854	168780.865	168880.142	0.094
No_Genre_Speechiness	338.970	168779.839	168779.848	168870.853	0.094
With_Genre	322.393	167327.793	167327.814	167468.452	0.138
Genre_Interaction	306.681	165869.012	165869.328	166423.371	0.182
Pop	322.907	41660.728	41660.765	41736.449	0.115
Rock	331.827	43532.405	43532.429	43594.707	0.044
HipHop	375.349	8312.993	8313.053	8344.467	0.025
Metal	223.664	21249.271	21249.308	21299.476	0.125
Electronic	315.678	43200.770	43200.811	43283.854	0.126
Classical	204.074	7559.464	7559.566	7601.537	0.131

1.6 Conclusions

1.6.1 Global model

Including genre improves model performance, with the adjusted R^2 accuracy increasing from approximately 0.09 to 0.18. Coefficient analysis shows that loudness and danceability have strong positive effects, whereas energy and instrumentalness have negative effects. ANOVA results confirm that loudness and danceability are the most important predictors. Despite this, overall explanatory power remains limited. Diagnostic checks indicate some assumption violations, but these are not critical given the large sample size.

1.6.2 Genre-specific models

Genre-specific models show both stable and varying effects. Danceability is consistently positive and liveness and valence consistently negative, while loudness positive in most genres indicating robust relationships. Other variables, including energy, instrumentalness, and duration, vary substantially, confirming that their effects depend on genre.

1.6.3 Limitations

The dataset appears artificially balanced across genres and does not include external factors such as artist popularity or marketing, limiting explanatory power. Some violations of model assumptions were observed, although these are mitigated by the large sample size. Additionally, grouping genres into broader categories may mask variation within subgenres. Tracks may also be associated with multiple genre labels in the original dataset; collapsing these to a single genre per track introduces a simplifying assumption that may not fully reflect the multi-genre nature of some songs.
